

Practical deep learning

Assignment 2

GitHub project:




Section 3.1.1 – Synthetic Data:

In this section we have generated random data, composed of 10,000 sequences, each containing $T = 50$ numbers.


For each of the sequences marked (x_1, \dots, x_{50}) , $x_i \in [0,1]$ and $\frac{1}{50} \sum x_i \approx \frac{1}{2}$

The data was split into 3 data sets: train, validation and test, with ratios of 60%, 20% and 20% accordingly.


Presented here are 3 examples of the random signals we generated:



signal Example 2



signal Example 3



Section 3.1.2

For this section, we implemented an LSTM Auto Encoder:

LSTM_AE(input_size, hidden_size, output_size, Labels = 1)

For this task:

input_size = output_size = 1 (dimension of each examples in a sequence – in this case each $x_i \in R$).

The hidden_size is a parameter for the grid search performed.

The AE is composed of:

1. Encoder = LSTM model with input dimension of 1 (each feature x_i is a number)

The output of the encoder is the context vector marked z, which is defined is the last hidden state, sized hidden_size (parameter of the network).

2. Decoder = LSTM model with input dimension of hidden size, which receives as input the context z repeated T (sequence length) times.

Output is of same size as input.

3. Fully Connected Layer – takes the activation of dimension (batch,T,hidden_size) and transforms it to the dimension (batch,T,hidden_size, Labels). In our case Labels =1 and we receive an output of the same dimension as the original input for the net.

Grid search

When training, we performed a grid search for the optimal parameters, training the net with chosen parameters on the train data, and testing it on the validation data every 5 epochs. for each parameters set we ran 300 epochs The optimizer used was Adam.

The range of parameters examined was:

Learning rate = [0.001,0.01]


Gradient clipping = [1,0(None)]

Hidden size = [30,60,120].

Finally, the optimal parameters chosen were:


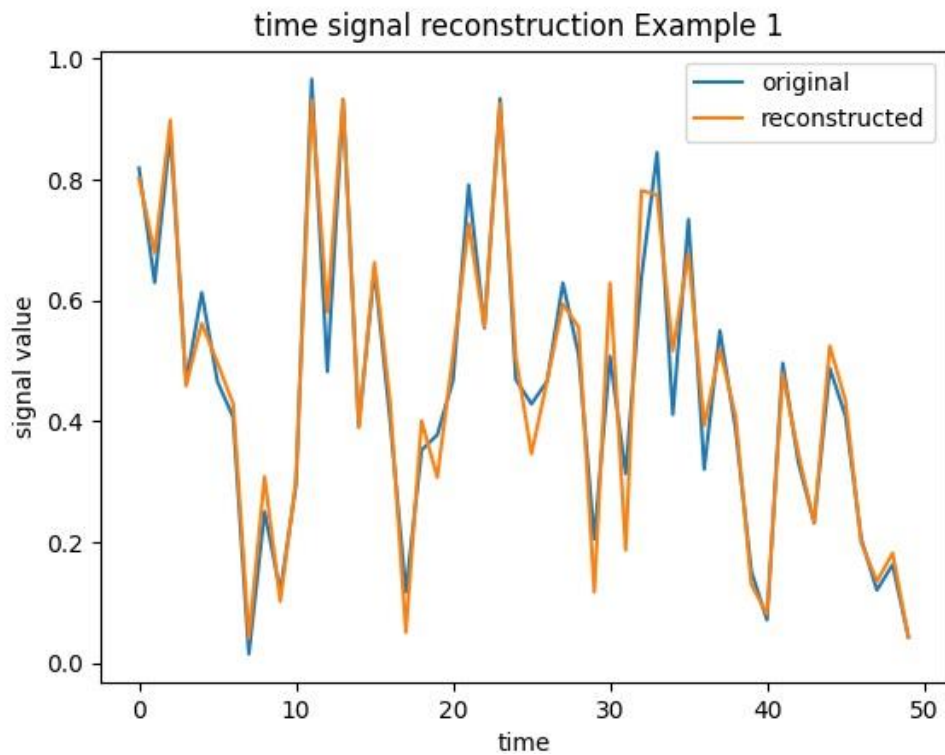
Learning rate = 0.001, gradient clipping = 1, Hidden size =60.

Below is the plot of the loss (MSE) calculated on the validation data each 5 epochs:



Finally, we ran the model trained with the chosen parameters on the test data, to receive the loss of 0.0004.

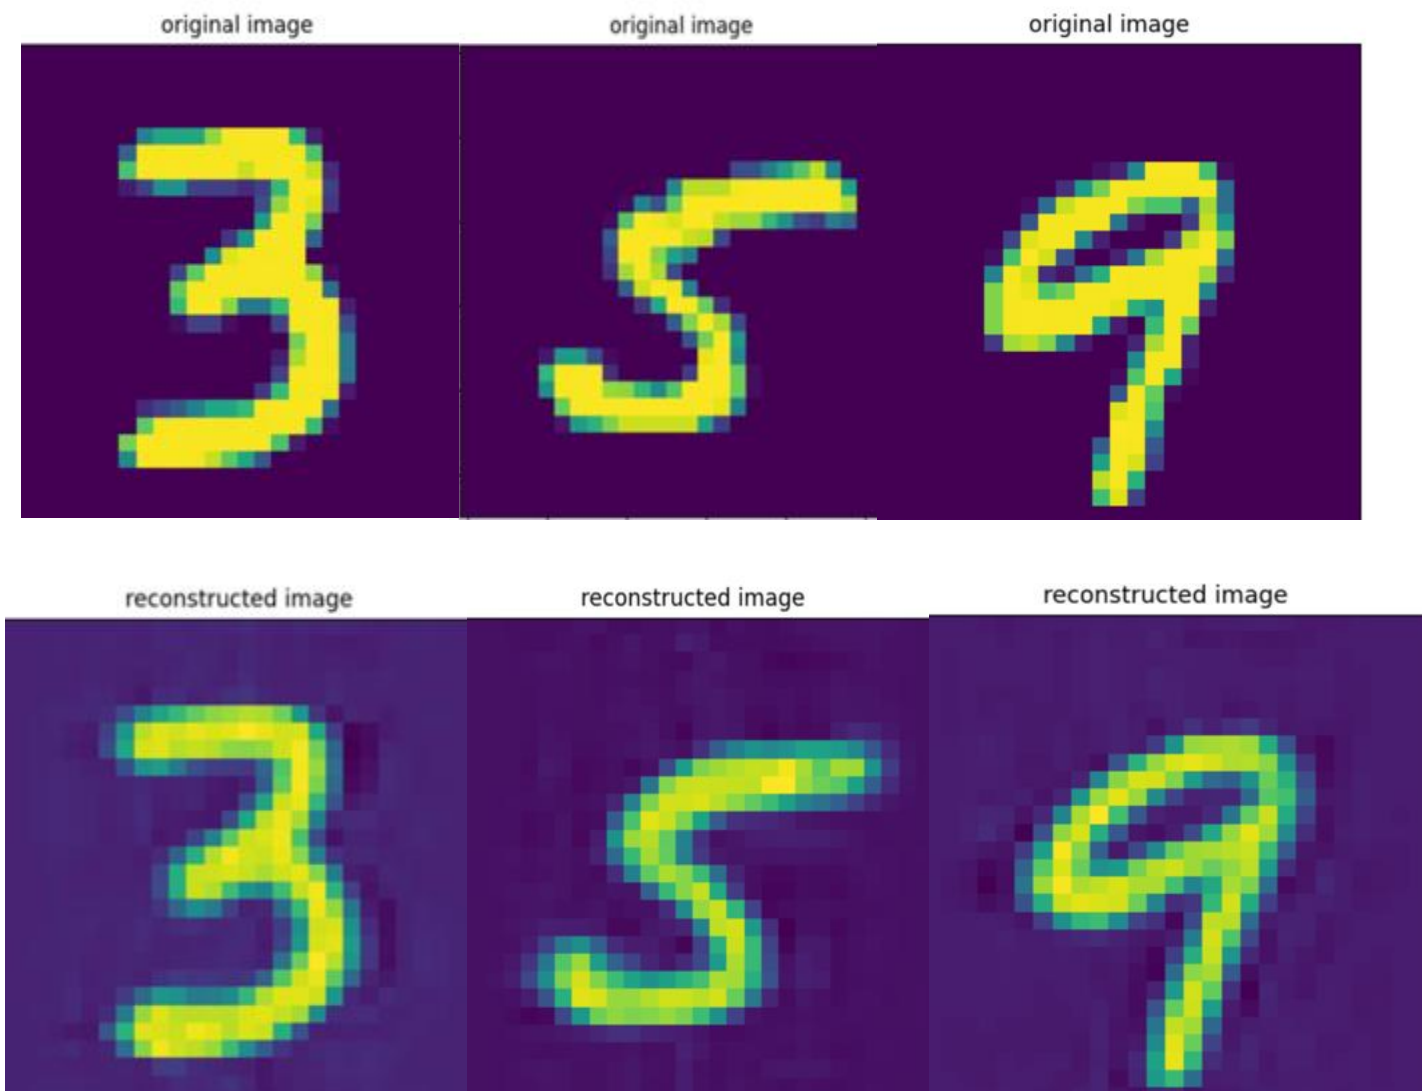
In the following figures, we present 2 examples of signals from the test data, reconstructed by our model:



Section 3.2- MNIST

During this section, we trained both a reconstruction and classification AE models on MNIST data. Due to computational limitations (very high training time), we used each image as a sequence of 28 rows (feeding the AE row by row instead of pixel by pixel). We used the chosen parameters from the grid search of the last section.

3.2.1: presented in the next 3 figure pairs, are the reconstruction of digits our model performed on the test data, post training:




3.2.2 – Classification:


During this section we performed classification on the images, meaning the last FC layer was changed to output a “probabilities” vector sized 10 for each input sequence, and we added to our MSE reconstruction loss a second Cross Entropy loss, making the loss for the entire network:

$$\text{Total loss} = \frac{MSE(\text{original } X, \text{reconstructed } X) + CE(\text{labels}, \text{predictions})}{2}$$

Training was very fast, reaching over 98% of accuracy in less than 10 epochs.

Below we present graphs for both validation loss and accuracy vs epochs:






Section 3.3 – S&P500

In this section, we performed reconstruction and classification on data of S&P500 stocks, working with the “high” (max) stock values per day.

Plotted below, we can see an example of the Amazon and Google high values of the stocks, between 01/14 – 01/18.



In order to normalize the data according to the requirements, we defined a `Normalizer()` class, performing the following normalization:

stock sequence rescaled = $\frac{\text{stock sequence} - \min(\text{stock sequence})}{\max(\text{stock sequence})}$, making the sequence values range in $[0,1]$.


then normalized stock sequence = $\frac{\text{stock sequence rescaled}}{2 * \text{mean}(\text{stock sequence rescaled})}$

setting the mean as $\frac{1}{2}$. After reconstruction, we will apply the opposite transformation to unnormalize the reconstructed stocks back to the original scale.


During training, we again used a validation data set to select optimal hyper parameters, finally choosing:

Hidden size = 120, learning rate = 0.001, gradient clipping = 1.


The following 3 graphs present reconstruction of 3 different stocks from our test data, after training the model with chosen hyper parameters:



TPR Stock reconstruction example



SBAC Stock reconstruction example



3.3.3: in this section, in order to perform classification we added a FC layer for prediction, acting on the decoded output in parallel to the reconstruction FC layer. For each stock high value x_1, \dots, x_{T-1} , we predicted $\hat{x}_2, \dots, \hat{x}_T$, meaning the value of the next day. For the loss of the network, we defined:

$$\text{Loss} = \frac{\text{MSE}(\text{stocks}, \text{reconstructed stocks}) + \text{MSE}(\text{predicted stocks}, \text{target stocks})}{2}$$

Below we present the plots of training and test loss vs epochs (3K epochs in total):

